



Linear discrimination

Jayanta Mukhopadhyay
Dept. of Computer Science and Engg.



Books

- Chapter 10 and 13 of “Introduction to Machine Learning” by Ethem Alpaydin.
- Chapter 5 of “Pattern Classification” by R.O. Duda, P. E. Hart and D. G. Stork



Discriminant functions

- Choose C_i if $g_i(\mathbf{x}) = \max_j g_j(\mathbf{x})$
- Linear function
 - $g_i(\mathbf{x} | \mathbf{w}_i, w_{i0}) = \mathbf{w}_i^T \mathbf{x} + w_{i0} = \sum_{j=1}^d w_{ij} x_j + w_{i0}$
 - Simple model, linear in form
 - $O(d)$ storage and time of computing $g(\cdot)$.
- Quadratic function
 - $g_i(\mathbf{x} | \mathbf{W}_i, \mathbf{w}_i, w_{i0}) = \mathbf{x}^T \mathbf{W}_i \mathbf{x} + \mathbf{w}_i^T \mathbf{x} + w_{i0}$
 - $O(d^2)$ storage and time of computing $g(\cdot)$.

Two classes

- One discriminant function sufficient

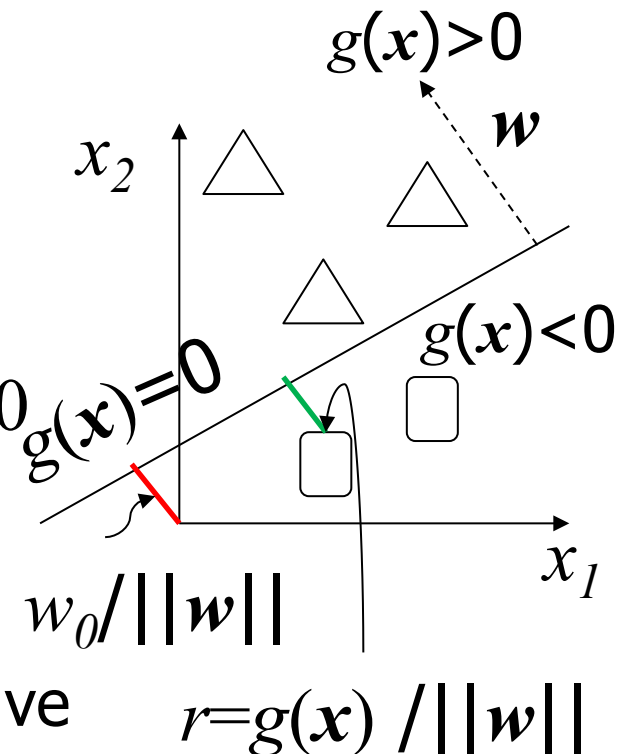
- $g(\mathbf{x}) = g_1(\mathbf{x} | \mathbf{w}_1, w_{10}) - g_2(\mathbf{x} | \mathbf{w}_2, w_{20})$
- $= \mathbf{w}_1^T \mathbf{x} + w_{10} - \mathbf{w}_2^T \mathbf{x} - w_{20}$
- $= (\mathbf{w}_1 - \mathbf{w}_2)^T \mathbf{x} - (w_{10} - w_{20})$
- $= \mathbf{w}^T \mathbf{x} + w_0$

- If $g(\mathbf{x}) > 0$ assign C_1 else C_2 .

- Hyper-plane dividing classes: $g(\mathbf{x}) = 0$

- Extend to more than 2 classes

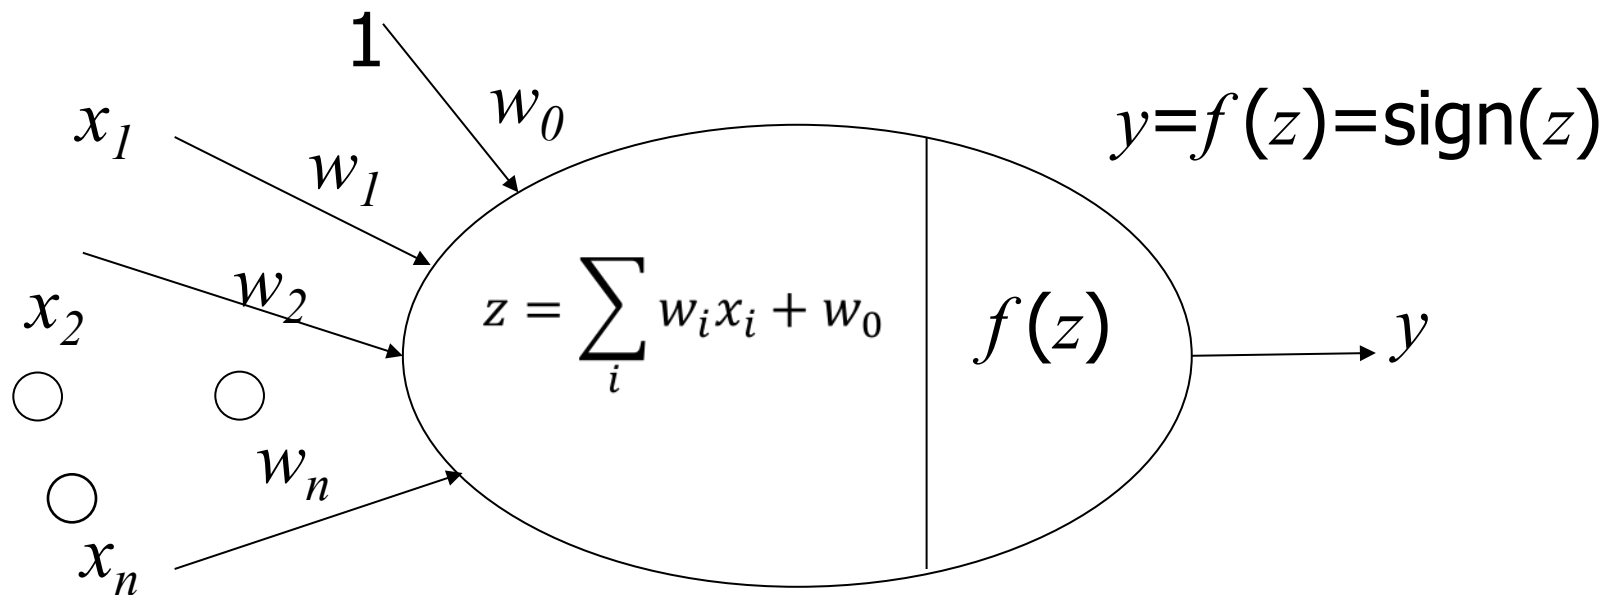
- Pairwise separation.
- Only samples of the class lies in the +ve half.






Perceptron classifier

- A linear classifier with a different perspective.



Augmented input and linear form




Given $\{y_i, X_i\}$, compute optimum W minimizing classification error.

Interpretation of $W^T X$

Consider the hyperplane $W^T X = 0$ separating two samples X and X' of classes 1 and 2.

w = unit vector
along W .

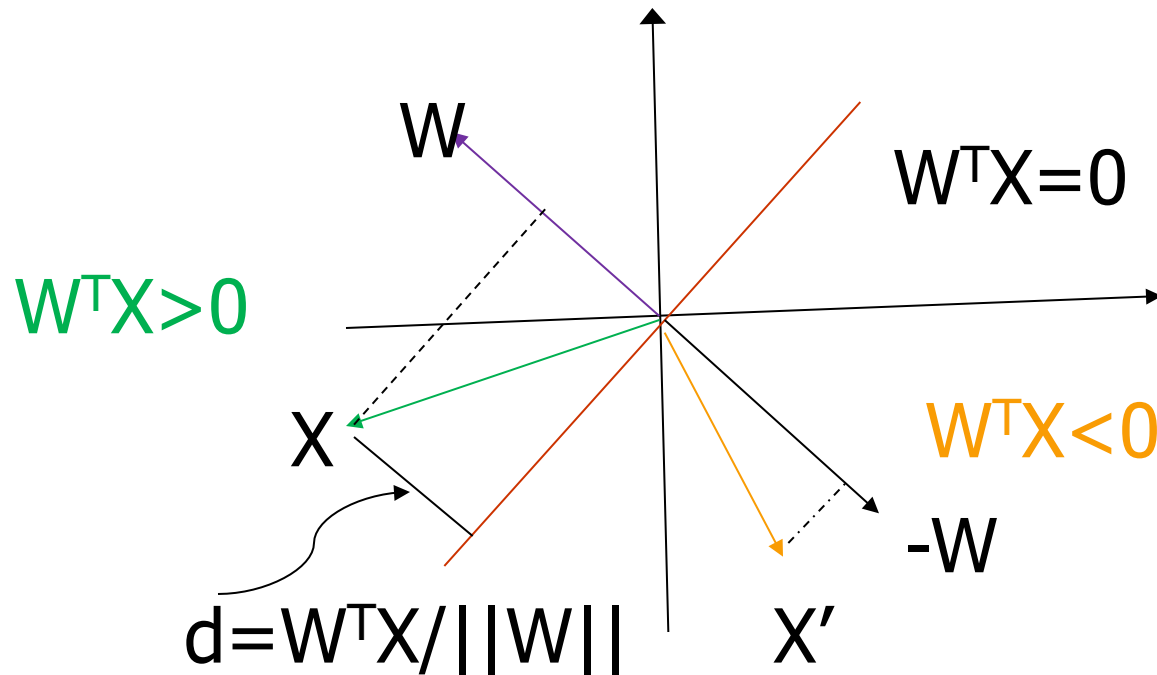


Distance of X from the hyperplane.

Linearly separable classes

To find a hyperplane separating data points of two classes.

If a solution exists, the classes are called linearly separable.

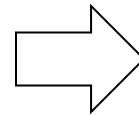


Distance of X from the hyperplane.

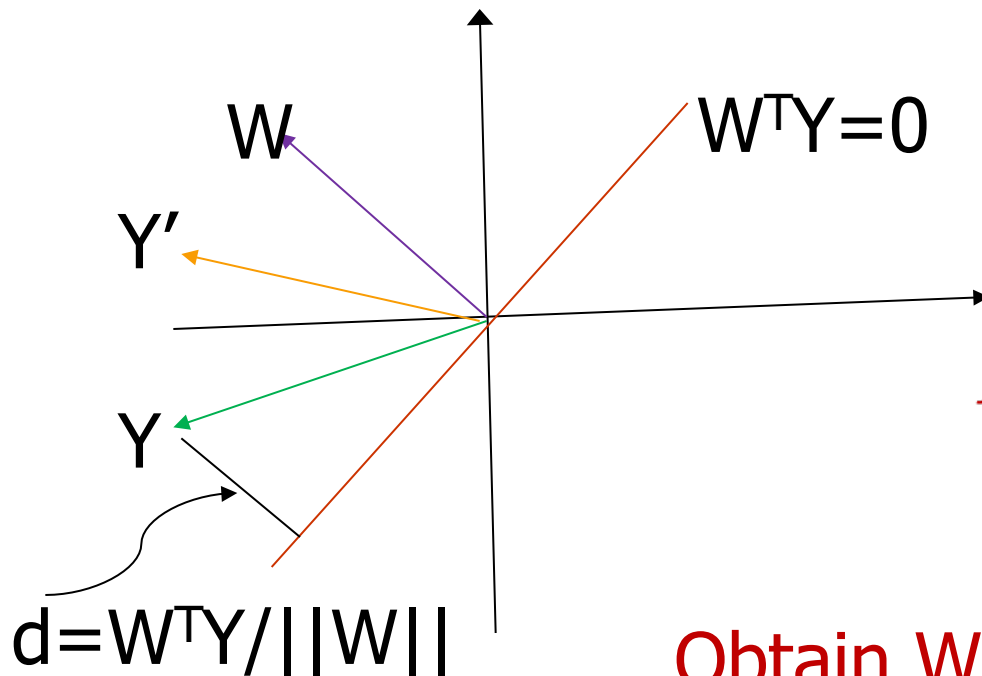
An error function

Data Normalization:

$Y = X$, if X in class 1 ($o = 1$).
 $= -X$, if X in class 2 ($o = -1$)



For correct classification,
 $W^T Y > 0$, for all Y .



Error function
(Perceptron Criterion):

$$J(W) = \sum_{Y \text{ misclassified}} -W^T Y$$

Always +ve

Obtain W which minimizes $J(W)$.



Gradient descent method for iterative optimization

- To obtain W which minimizes $J(W)$.
- Start with an initial vector $W^{(0)}$.
- Compute the gradient vector $\nabla J(W^{(0)})$
- Move closer to minimum by updating W .

$$W^{(i)} = W^{(i-1)} - \eta(i) \nabla J(W)$$

Positive scale factor
(learning rate)



Iterative gradient descent Optimization

Data Normalization:

$Y=X$, if X in class 1 ($o=1$).
 $=-X$, if X in class 2 ($o=-1$)

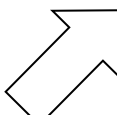
$$J_p(W) = \sum_{Y \text{ misclassified.}} -W^T Y$$

Iterative Optimization using gradient descent

1. Start with $W^{(0)}$.

2. Update W

$$W^{(i)} = W^{(i-1)} + \eta(i) \sum_{Y \text{ misclassified}} Y$$


$$W^{(i)} = W^{(i-1)} - \eta(i) \nabla J_p(W)$$

May be taken as a constant.

3. Continue step 2 till converges.

Other forms of the error function

■ There could be other forms of the criterion function.

- $J_p(W)$: not continuous
- $J_q(W)$: continuous.
 - Very smooth in boundary.
 - May get stuck there.
 - Value dominated by long Y 's.

Gradient computation:

$$\nabla J_r(W) = \sum_{W^T Y \leq b} \frac{Y(W^T Y - b)}{\|Y\|^2}$$

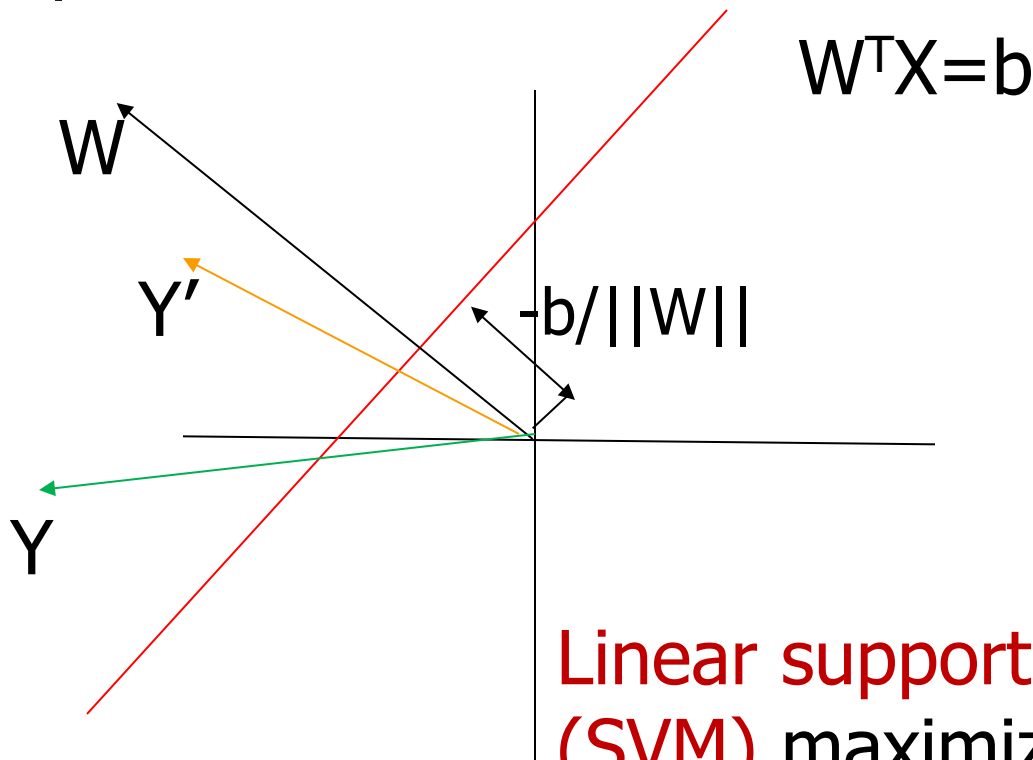
$$J_q(W) = \sum_{Y \text{ misclassified.}} (W^T Y)^2$$

Another error function
(Relaxation criterion)

$$J_r(W) = \frac{1}{2} \sum_{\substack{Y \text{ misclassified} \\ W^T Y \leq b}} \frac{(W^T Y - b)^2}{\|Y\|^2}$$

↑
Stronger linear separability

More stringent criteria of linear separability



Linear support vector machines (SVM) maximize this margin of separation between two linearly separable data points of classes.



The algorithm (Batch relaxation with margin)

- Initialize W to $W^{(0)}$.
- Iterate till convergence
- Compute the set M of misclassified samples (with margin b), so that
 - $M = \{Y | W^T Y \leq b\}$
- Compute gradient.
- Update W .

$$\nabla J_r(W) = \sum_{W^T Y \leq b} \frac{Y(W^T Y - b)}{\|Y\|^2}$$

$$W^{(i)} = W^{(i-1)} - \eta(i) \nabla J_r(W^{(i-1)})$$



Single sample relaxation with margin

- Initialize W to $W^{(0)}$.
- Perform the update of W by considering samples one by one in every iteration.
- Consider an i th sample Y_i at k th iteration.
- If $(W^T Y_i \leq b)$
 - Update W .
$$W^{(k)} = W^{(k-1)} + \eta(k) \frac{b - W^T Y_i}{\|Y_i\|^2} Y_i$$
- Stop when very little change in updates at the end of an iteration.



Support Vector Machine (SVM)

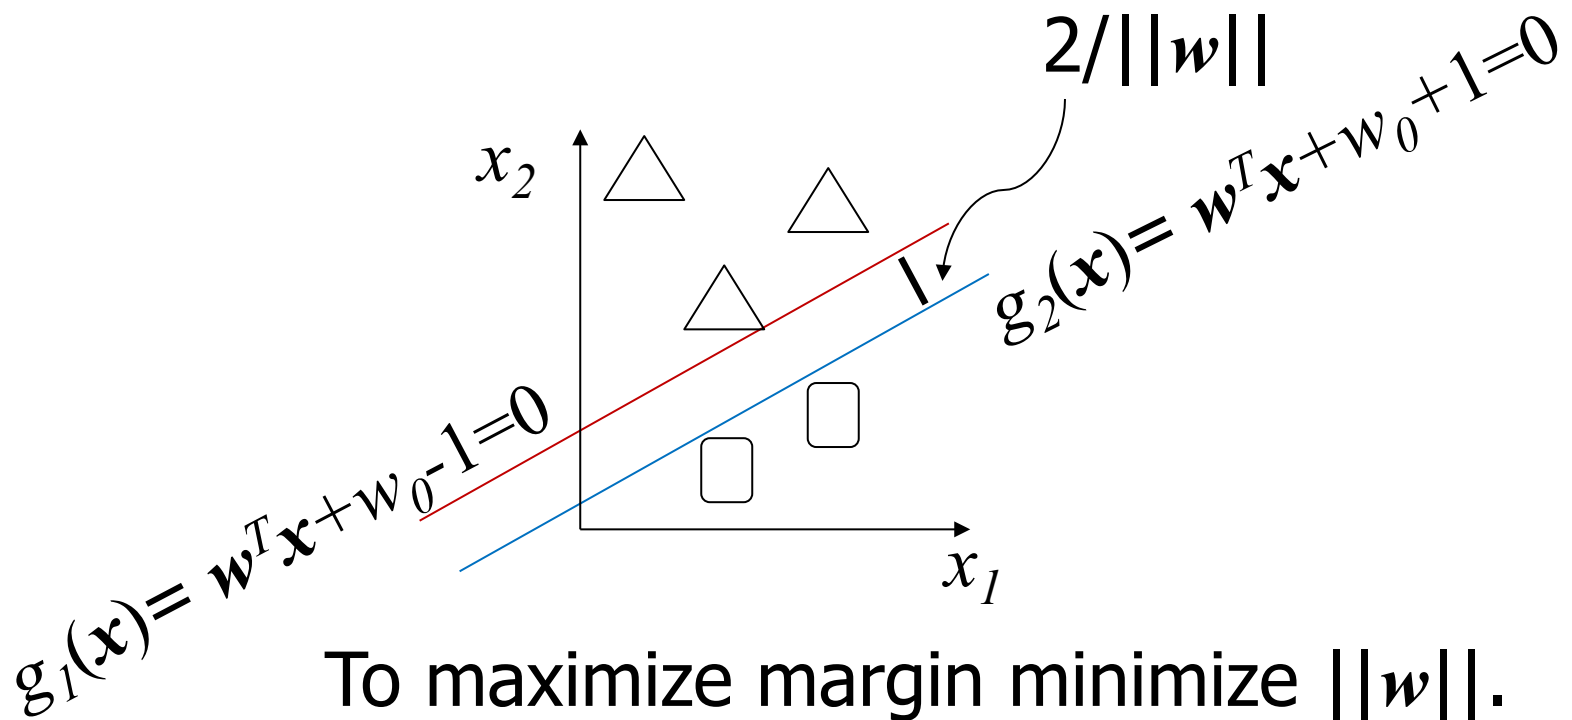
- A linear discriminant classifier.
- Uses Vapnik's principle:
 - to never solve a more complex problem as a first step before the actual problem.
 - Classification: Sufficient to compute class boundaries (where $P(C_1|\mathbf{x})=P(C_2|\mathbf{x})$) without computing class distributions $P(C_i|\mathbf{x})$, etc.
 - Outlier detection: Compute boundaries separating those \mathbf{x} having low $P(\mathbf{x})$.
- After training the weight vector can be written in terms of training samples lying in class boundaries.
 - Support vectors.



Two class problem

- $X = \{\mathbf{x}^t, r^t\}, t = 1, 2, \dots, N$
 - \mathbf{x}^t in R^d , r^t in $\{+1, -1\}$.
- To compute \mathbf{w} and w_0 such that for all t
 - $\mathbf{w}^T \mathbf{x}^t + w_0 > +1$ if $r^t = +1$
 - $\mathbf{w}^T \mathbf{x}^t + w_0 < -1$ if $r^t = -1$
- Rewritten as:
 - $r^t(\mathbf{w}^T \mathbf{x}^t + w_0) \geq +1$ for all t
 - Note: it is harder than $r^t(\mathbf{w}^T \mathbf{x}^t + w_0) \geq 0$
 - A margin left between zones of two classes.

Margin between classes



To minimize $||\mathbf{w}||^2/2$ subject to $r^t(\mathbf{w}^T \mathbf{x}^t + w_0) \geq +1$ for all t



Optimization problem

- Constrained optimization problem

- To minimize $\|\mathbf{w}\|^2/2$

- subject to $r^t(\mathbf{w}^T \mathbf{x}^t + w_0) \geq +1$ for all t

- Unconstrained problem:

$$L_p = \frac{1}{2} \|\mathbf{w}\|^2 - \sum_{t=1}^N \alpha^t [r^t(\mathbf{w}^T \mathbf{x}^t + w_0) - 1]$$

Lagrange multipliers

$$L_p = \frac{1}{2} \|\mathbf{w}\|^2 - \sum_{t=1}^N \alpha^t r^t(\mathbf{w}^T \mathbf{x}^t + w_0) + \sum_{t=1}^N \alpha^t$$

- To be minimized w.r.t \mathbf{w} and w_0 and maximized w.r.t. Lagrange multipliers.



Convex quadratic optimization problem

- Unconstrained problem:

$$L_p = \frac{1}{2} \|\mathbf{w}\|^2 - \sum_{t=1}^N \alpha^t r^t (\mathbf{w}^T \mathbf{x}^t + w_0) + \sum_{t=1}^N \alpha^t$$

Lagrange multipliers

- Convex objective function and linear constraints.

- Dual problem

- To be maximized w.r.t. Lagrange multipliers (>0)
subject to that gradients w.r.t \mathbf{w} and w_0 should be 0.

$$\frac{\partial L_p}{\partial \mathbf{w}} = 0 \Rightarrow \mathbf{w} = \sum_t \alpha^t r^t \mathbf{x}^t \qquad \frac{\partial L_p}{\partial w_0} = 0 \Rightarrow \sum_t \alpha^t r^t = 0$$



Dual optimization problem

- Primary problem:

$$L_p = \frac{1}{2} \|\mathbf{w}\|^2 - \sum_{t=1}^N \alpha^t r^t (\mathbf{w}^T \mathbf{x}^t + w_0) + \sum_{t=1}^N \alpha^t$$

- Dual problem derived by applying following conditions in the primary objective function.

$$\frac{\partial L_p}{\partial \mathbf{w}} = 0 \Rightarrow \mathbf{w} = \sum_t \alpha^t r^t \mathbf{x}^t \qquad \frac{\partial L_p}{\partial w_0} = 0 \Rightarrow \sum_t \alpha^t r^t = 0$$

$$L_d = \frac{1}{2} \mathbf{w}^T \mathbf{w} - \mathbf{w}^T \overbrace{\sum_t \alpha^t r^t \mathbf{x}^t}^{\mathbf{w}} - w_0 \underbrace{\sum_t \alpha^t r^t}_0 + \sum_t \alpha^t$$

\Downarrow

$$L_d = -\frac{1}{2} \mathbf{w}^T \mathbf{w} + \sum_t \alpha^t$$



Dual optimization problem

- Dual problem:

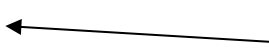
$$L_d = -\frac{1}{2} \mathbf{w}^T \mathbf{w} + \sum_t \alpha^t$$

$$\frac{\partial L_p}{\partial \mathbf{w}} = 0 \Rightarrow \mathbf{w} = \sum_t \alpha^t r^t \mathbf{x}^t$$

$$\frac{\partial L_p}{\partial w_0} = 0 \Rightarrow \sum_t \alpha^t r^t = 0$$
$$\alpha^t \geq 0$$

- Expand \mathbf{w} from the condition.

$$L_d = -\frac{1}{2} \sum_t \sum_s \alpha^t \alpha^s r^t r^s (\mathbf{x}^t)^T \mathbf{x}^s + \sum_t \alpha^t$$

- To maximize L_d w.r.t. α^t  Most of them
- Apply quadratic optimization technique: will be 0.
 - $O(N^3)$ time and $O(N^2)$ space complexity.



Solution

- Dual problem:

$$L_d = -\frac{1}{2} \sum_t \sum_s \alpha^t \alpha^s r^t r^s (\mathbf{x}^t)^T \mathbf{x}^s + \sum_t \alpha^t$$

$$\frac{\partial L_p}{\partial \mathbf{w}} = 0 \Rightarrow \mathbf{w} = \sum_t \alpha^t r^t \mathbf{x}^t$$

$$\frac{\partial L_p}{\partial w_0} = 0 \Rightarrow \sum_t \alpha^t r^t = 0$$

$$\alpha^t \geq 0$$

- Apply quadratic optimization technique:

- $O(N^3)$ time and $O(N^2)$ space complexity.

- Most of α^t will be zero.

- Samples with positive (non-zero) α^t are support vectors.

- Provide \mathbf{w} as a linear combination of input samples (Condition 1).

- w_0 is obtained from **any** of the support vector which lies in the boundary.

- $r^t(\mathbf{w}^T \mathbf{x}^t + w_0) = +1 \Rightarrow w_0 = r^t - \mathbf{w}^T \mathbf{x}^t$ (For numerical stability take **avg.**).

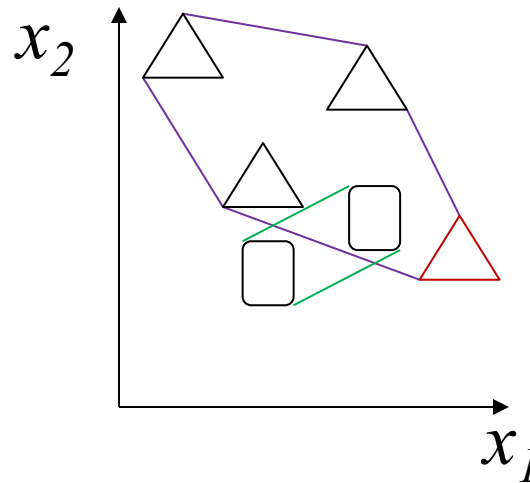


SVM- Testing

- Check only the sign of discriminant value.
 - Margin not enforced.
- Only support vectors decide class boundaries.
 - Other samples do not influence the classifier.

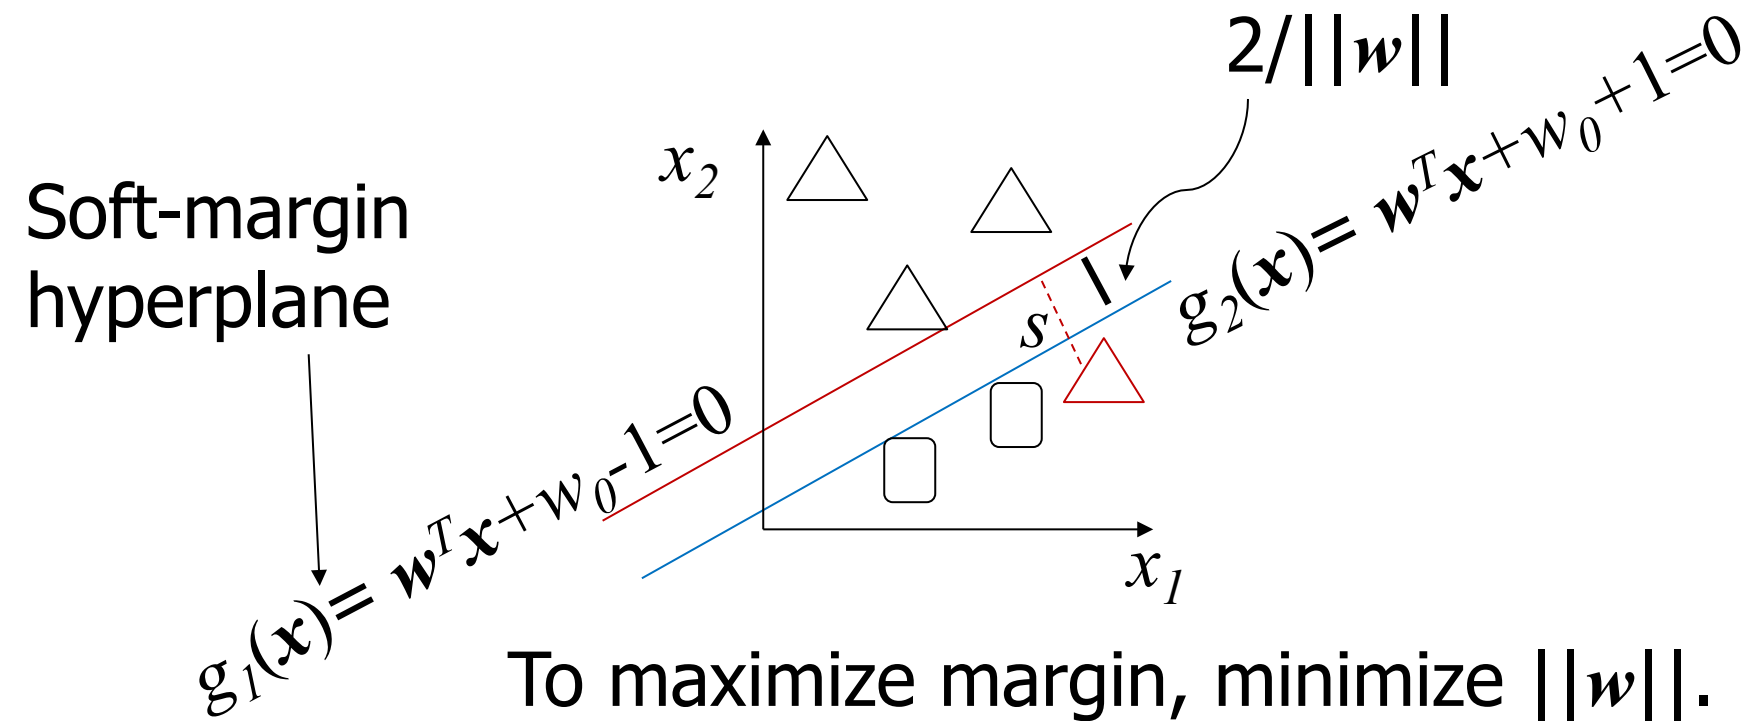
The non-separable case: Soft margin hyperplane

- Classes may not be linearly separable.



- Use of slack variable, $\{s^t\}$, $t=1,2,..N$

Slack variable to define soft margin



To minimize $\|\mathbf{w}\|^2/2$ subject to $r^t(\mathbf{w}^T \mathbf{x}^t + w_0) \geq +1 - s^t$ for all t



Constraints with Soft margin hyperplanes

- Use of slack variable, $\{s^t\}$, $t=1,2,..N$ to define constraints.
 - $r^t(\mathbf{w}^T \mathbf{x}^t + w_0) \geq 1 - s^t$ for all t
 - $0 < s^t < 1$, \mathbf{x}^t correctly classified.
 - If $s^t \geq 1$, \mathbf{x}^t misclassified.
 - $\#[s^t > 1]$: Number of misclassified points.
 - $\#[s^t > 0]$: Number of non-separable points.
 - Soft error = $\sum_t s^t$



Optimization problem

- Add penalty term for soft error to define the objective function for minimization.

- $L_p = \|\mathbf{w}\|^2/2 + C \sum_t s^t$
 - subject to $r^t(\mathbf{w}^T \mathbf{x}^t + w_0) \geq 1 - s^t$ for all t
 - C is the penalty factor.

$$L_p = \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_t s^t - \sum_{t=1}^N \alpha^t [r^t(\mathbf{w}^T \mathbf{x}^t + w_0) - 1 + s^t] - \sum_t \mu^t s^t$$

- μ^t are the new Lagrange parameters to guarantee that $s^t > 0$.



Optimization problem

- Primary problem:

$$L_p = \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_t s^t - \sum_{t=1}^N \alpha^t [r^t (\mathbf{w}^T \mathbf{x}^t + w_0) - 1 + s^t] - \sum_t \mu^t s^t$$

$$\frac{\partial L_p}{\partial \mathbf{w}} = 0 \Rightarrow \mathbf{w} = \sum_t \alpha^t r^t \mathbf{x}^t \quad \frac{\partial L_p}{\partial w_0} = 0 \Rightarrow \sum_t \alpha^t r^t = 0$$

$$\frac{\partial L_p}{\partial s^t} = C - \alpha^t - s^t = 0 \Rightarrow 0 \leq \alpha^t \leq C \text{ as } s^t > 0$$

- Dual problem: $L_d = -\frac{1}{2} \sum_t \sum_s \alpha^t \alpha^s r^t r^s (\mathbf{x}^t)^T \mathbf{x}^s + \sum_t \alpha^t$
 - Subject to: $\sum_t \alpha^t r^t = 0$ and $0 \leq \alpha^t \leq C$

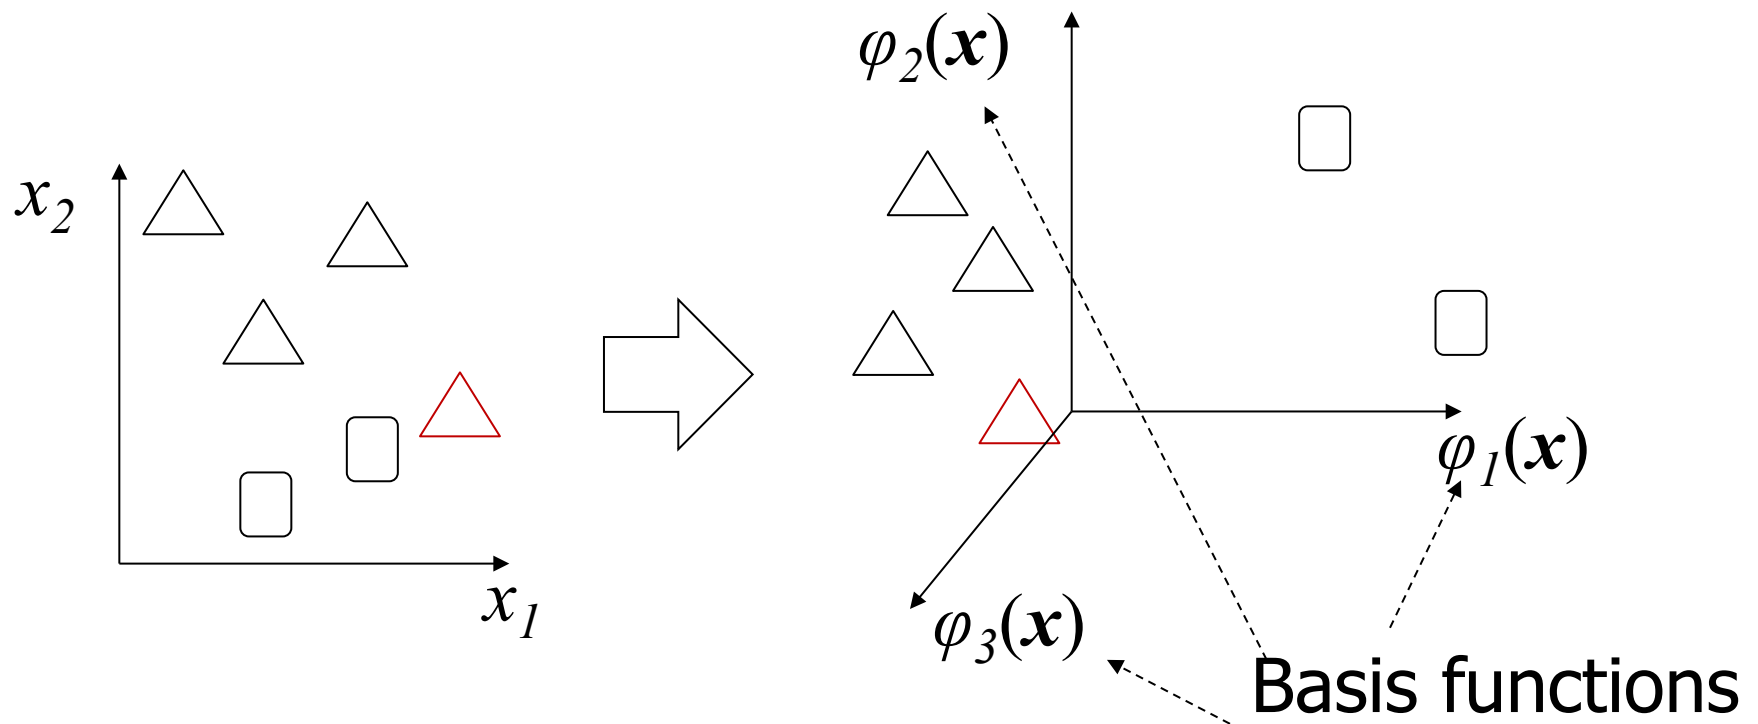


The solution

- The same quadratic optimization technique to be used.
- The support vectors have $\alpha^t > 0$
- Out of them whose values are less than C are used for deriving w_0 .
 - $w_0 = r^t - w^T x^t$
 - Take average.

Projecting to higher dimensional space

- May make them linearly separable!





Solving by projecting to a high-dimensional space.

- $\mathbf{z} = \varphi(\mathbf{x})$, where $z_j = \varphi_j(\mathbf{x})$, $j=1,2,\dots,k$
- $g(\mathbf{z}) = \mathbf{w}^T \mathbf{z}$
 - Assume $z_1=1$ (for taking care of the constant term w_0 as used previously).
- $g(\mathbf{x}) = \sum_j w_j \varphi_j(\mathbf{x})$
- No guarantee that the classes are linearly separable in the space of basis functions.
- Similar optimization problem:
 - $L_p = ||\mathbf{w}'||^2/2 + C \sum_t s^t$
 - subject to $r^t \mathbf{w}'^T \varphi(\mathbf{x}^t) \geq 1 - s^t$ for all t
 - C is the penalty factor.



Primal-Dual problems

- Primal problem:
- To minimize w.r.t \mathbf{w}

$$L_p = \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_t s^t - \sum_{t=1}^N \alpha^t [r^t (\mathbf{w}^T \boldsymbol{\varphi}(\mathbf{x}^t) + w_0) - 1 + s^t] - \sum_t \mu^t s^t$$

$$\frac{\partial L_p}{\partial \mathbf{w}} = 0 \Rightarrow \mathbf{w} = \sum_t \alpha^t r^t \boldsymbol{\varphi}(\mathbf{x}^t) \quad \frac{\partial L_p}{\partial s^t} = C - \alpha^t - s^t = 0 \quad \sum_t \alpha^t r^t = 0$$

- Dual problem

$$L_d = -\frac{1}{2} \sum_t \sum_s \alpha^t \alpha^s r^t r^s \overbrace{(\boldsymbol{\varphi}(\mathbf{x}^t))^T \boldsymbol{\varphi}(\mathbf{x}^s)}^{K(\mathbf{x}^t, \mathbf{x}^s)} + \sum_t \alpha^t$$

- Subject to $\sum_t \alpha^t r^t = 0$ and $0 \leq \alpha^t \leq C$

$K(\mathbf{x}^t, \mathbf{x}^s)$ ← Kernel function



Kernel machines

- Discriminant function

$$g(\mathbf{x}) = \mathbf{w}^T \boldsymbol{\varphi}(\mathbf{x}) = \sum_t \alpha^t r^t \boldsymbol{\varphi}(\mathbf{x}^t)^T \boldsymbol{\varphi}(\mathbf{x})$$



$$g(\mathbf{x}) = \sum_t \alpha^t r^t K(\mathbf{x}^t, \mathbf{x})$$

A real symmetric $n \times n$ matrix M is +ve semidefinite iff $\mathbf{z}^T M \mathbf{z} \geq 0$ for any non-zero \mathbf{z} in R^n .

- No need to compute with basis functions and also performing dot products with \mathbf{z} .
- Gram matrix: The matrix of kernel values \mathbf{K} , where $K_{t,s} = K(\mathbf{x}^t, \mathbf{x}^s)$
 - Should be symmetric and +ve semidefinite.
 - To be provided.



Vectorial kernel functions

- Polynomials of degree q
 - $K(\mathbf{x}^t, \mathbf{x}) = (\mathbf{x}^T \mathbf{x}^t + 1)^q$
- Radial basis functions
 - $K(\mathbf{x}^t, \mathbf{x}) = \exp[-\|\mathbf{x} - \mathbf{x}^t\|^2 / (2s^2)]$
- Mahalanobis kernel function
 - $K(\mathbf{x}^t, \mathbf{x}) = \exp[-(\mathbf{x} - \mathbf{x}^t)^T S^{-1} (\mathbf{x} - \mathbf{x}^t) / 2]$
- Distance function based
 - $K(\mathbf{x}^t, \mathbf{x}) = \exp[-D(\mathbf{x}, \mathbf{x}^t) / (2s^2)]$
- Sigmoidal function: $K(\mathbf{x}^t, \mathbf{x}) = \tanh(2\mathbf{x}^T \mathbf{x}^t + 1)$

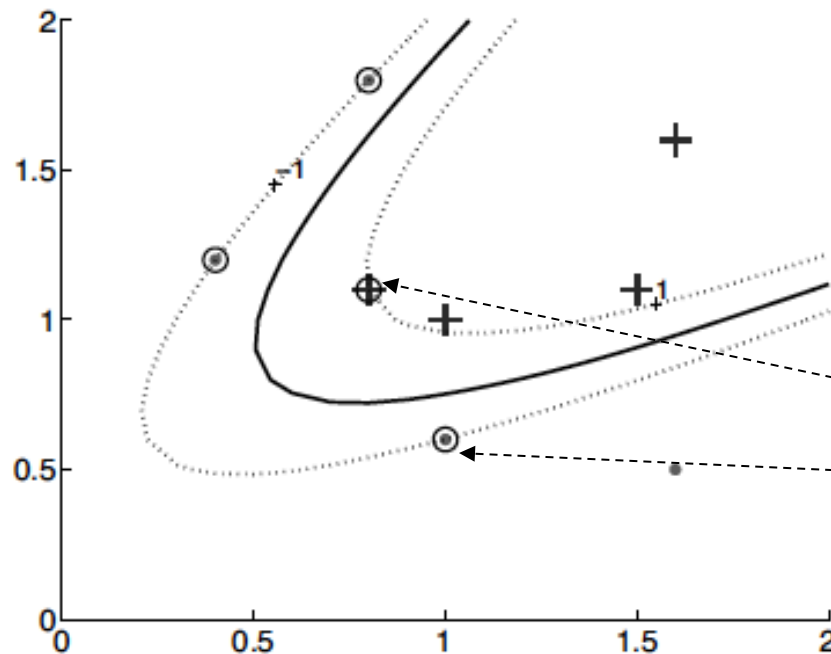
A typical example:

- Decision Boundaries of Quadratic Kernel in 2D:

- $K(x,y)=(x^T y+1)^2$

$$\varphi(\mathbf{x}) = \begin{bmatrix} 1 \\ \sqrt{2}x_1 \\ \sqrt{2}x_2 \\ x_1^2 \\ x_2^2 \end{bmatrix}$$

No need to
compute $\varphi(\mathbf{x})$.



Circled
instances are
support
vectors.



Defining kernels

- May be defined between a pair of objects flexibly.
 - without using any closed functional form.
- A few examples
 - Number of shared words of two documents.
 - Edit distance between two strings.
 - Number of shared paths between two graphs.
 - Empirical definition of a kernel matrix on training samples.
- The same principle applicable for designing SVM for classifying such objects.



SVM: Summary

- SVM provides maximum margin based linear discrimination for two linearly separable classes.
- Generalizes to non-separable classes by using slack variables.
- Use of basis functions to map non-separable classes to separable in a higher dimensional space.
 - Computation becomes simple and efficient with kernel functions.



Parametric discrimination revisited

- Class densities $P(\mathbf{x}|C_i)$ Gaussian sharing a common cov. Matrix: Σ , and $\boldsymbol{\mu}_i: E(\mathbf{x}|C_i)$.
 - The discriminant function is linear
 - $g_i(\mathbf{x}) = \mathbf{w}_i^T \mathbf{x} + w_{i0}$
 - $\mathbf{w}_i = \Sigma^{-1} \boldsymbol{\mu}_i$
 - $w_{i0} = -(\boldsymbol{\mu}_i^T \Sigma^{-1} \boldsymbol{\mu}_i)/2 + \log P(C_i)$
 - Two classes: Let $P(C_1|\mathbf{x}) = y$, hence $P(C_2|\mathbf{x}) = 1 - y$
 - Choose C_1 if $y > 0.5 \Leftrightarrow y/(1-y) > 1 \Leftrightarrow \log(y/(1-y)) > 0$
 - Else Choose C_2 .
- logit(y) or log odds of y .



Parametric discrimination revisited

- Two classes: Let $P(C_1|\mathbf{x})=y$, hence $P(C_2|\mathbf{x})=1-y$
 - Choose C_1 if $\text{logit}(y)(=\log(y/1-y))>0$, else Choose C_2 .
- For two normal classes sharing a common covariance matrix, the log odds linear.
- $\text{logit}(P(C_1|\mathbf{x}))=\mathbf{w}^T\mathbf{x}+w_0$
 - $\mathbf{w}=\Sigma^{-1}(\boldsymbol{\mu}_1-\boldsymbol{\mu}_2)$
 - $w_0=-((\boldsymbol{\mu}_1-\boldsymbol{\mu}_2)^T\Sigma^{-1}(\boldsymbol{\mu}_1-\boldsymbol{\mu}_2))/2 + \log P(C_1)/P(C_2)$
- The inverse of logit is the logistic function, also called *sigmoid function*.
- $P(C_1|\mathbf{x})=\text{logit}^{-1}(\mathbf{w}^T\mathbf{x}+w_0) = 1/(1+\exp(-(\mathbf{w}^T\mathbf{x}+w_0)))$



Parametric two class classification using discriminant

- Estimate parameters, Σ , μ_1 , and μ_2 .
- Compute coefficients of $g(\mathbf{x})$: \mathbf{w} , and w_0 .
- During testing:
 - Calculate $g(\mathbf{x})$.
 - Assign C_1 if $g(\mathbf{x}) > 0$, else C_2 .
- OR
 - Calculate $y = 1 / (1 + \exp(-(\mathbf{w}^T \mathbf{x} + w_0)))$
 - Assign C_1 if $y > 0.5$ else C_2 .



Logistic discrimination of two classes

- Ratio of class densities modeled: $P(\mathbf{x}|C_1)/P(\mathbf{x}|C_2)$
- Assume log likelihood ratio is linear
 - true for normal density functions.
- $\log(P(\mathbf{x}|C_1)/P(\mathbf{x}|C_2)) = \mathbf{w}^T \mathbf{x} + w_0'$
- ➔ $\text{logit}(P(C_1|\mathbf{x})) = \log(P(C_1|\mathbf{x})/P(C_2|\mathbf{x}))$
 $= \log(P(\mathbf{x}|C_1)/P(\mathbf{x}|C_2)) + \log(P(C_1)/P(C_2))$
 $= \mathbf{w}^T \mathbf{x} + w_0 \quad (\text{when } w_0 = w_0' + \log(P(C_1)/P(C_2)))$
- $y = P(C_1|\mathbf{x}) = 1/(1 + \exp(-(\mathbf{w}^T \mathbf{x} + w_0)))$



Learning weights of logit functions.

- Data: $X = \{\mathbf{x}^t, r^t\}, t=1, 2, \dots, N$
 - $r^t = 1$ for C_1 , and 0 for C_2 .
- Let $y = P(r^t = 1 | \mathbf{x}) \sim \text{Bernoulli}(y)$.
 - Directly modeling likelihood of class assignment
 - instead of likelihood of data given classes as in the parametric approach.

$$l(\mathbf{w}, w_0 | X) = \prod_{t=1}^N (y^t)^{r^t} (1 - y^t)^{(1-r^t)}$$

- To minimize $E = -\log(l)$ (Maximization of log likelihood)

$$E = -\left(\sum_t (r^t \log y^t + (1 - r^t) \log(1 - y^t))\right)$$

- Use gradient descent technique to iterate on weights.

$$y = 1/(1 + \exp(-a))$$



Gradient descent technique

- $y = \text{sigmoid}(a) \Rightarrow dy/da = y \cdot (1 - y)$

$$\begin{aligned} \frac{\partial E}{\partial w_j} &= - \sum_t \left(\frac{r^t}{y^t} - \frac{1 - r^t}{1 - y^t} \right) y^t (1 - y^t) x_j^t \\ &= - \sum_t (r^t - y^t) x_j^t \end{aligned}$$

$$\frac{\partial E}{\partial w_0} = - \sum_t (r^t - y^t)$$

- Update of weights at i th iteration.

$$w_j^{(i)} = w_j^{(i-1)} - \eta \frac{\partial E}{\partial w_j}$$



Algorithm (Learning weights)

1. Assume initial \mathbf{w} , and w_0 .
2. Compute $y = \text{sigmoid}(\mathbf{w}^T \mathbf{x} + w_0)$
3. Compute gradients.
4. Update \mathbf{w} , and w_0 .
5. Continue steps 2 to 4 till convergence.

Extended to multi-class
problem by modeling
 $P(C_i|\mathbf{x})$ by softmax(.)
function, and multinomial
distr.

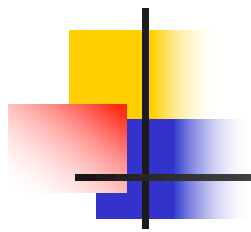
$$y_i = P(C_i|\mathbf{x}) = \frac{\exp(\mathbf{w}_i^T \mathbf{x} + w_{i0})}{\sum_j \exp(\mathbf{w}_j^T \mathbf{x} + w_{j0})}$$

$$l(\{\mathbf{w}_i, w_{i0}\}|X) = \prod_t \prod_i (y_i^t)^{r_i^t}$$



Summary

- Discriminant functions could be explained in the context of Bayesian inference.
- Could be explained by geometry.
- Weights of the function to be learned by minimizing an objective function (error due to miss-classification).
 - Gradient descent method.
 - Stochastic gradient descent.
- Linear SVM: optimally separable hyperplane.
- Logistic discrimination: regresses posterior directly from labelled data.



Thank you!